

PostgreSQL and MariaDB under TPC-C: SSD vs. RAID HDD

Gustaw Blachowski

Faculty of Mathematics, Informatics and Mechanics, University of Warsaw

Supervisor: prof. dr hab. Krzysztof Stencel

Abstract

In this project I compared how the storage drive (SSD vs. a RAID array of HDDs) and the database engine (PostgreSQL vs. MariaDB) affect performance on the TPC-C benchmark. In my first round of tests, run with `SERIALIZABLE` isolation, most HDD transactions were aborted, and attempts to reproduce the results ran into PostgreSQL's lock table filling up, so I switched to `READ COMMITTED`. PostgreSQL on SSD was then about 3x faster than on HDD at moderate load, but the HDD array nearly caught up at 70 concurrent terminals. The SSD's latency stayed lower at every load. MariaDB was about 3x slower than PostgreSQL on SSD and over 30x slower on HDD, but this is not a fair comparison of the engines. Both have the same small internal cache by default, but PostgreSQL also uses the OS page cache, which can grow into all free memory on the server, while MariaDB uses direct I/O and only has its buffer pool, which I couldn't resize.

1. Introduction

The usual assumption is that a database on an SSD will be much faster than one on HDDs. In practice, the difference depends on how much data is served from memory, how the database engine does I/O, and how much time goes into locking rather than disk access. In this project I wanted to check how big the difference really is on a realistic workload, and whether the choice of database engine matters more than the choice of drive.

In this report I describe:

1. A comparison of PostgreSQL on an SSD and on a RAID array of HDDs, with 1 to 70 concurrent terminals.
2. My first round of tests under `SERIALIZABLE` isolation, and the lock problems I ran into when trying to reproduce it (Section 4.1).
3. An attempt to compare PostgreSQL and MariaDB on the same hardware and workload, and why their different caching (page cache vs. direct I/O) made a fair comparison impossible in my setup.

2. Background

TPC-C [1] simulates an order-processing system with five transaction types (New-Order, Payment, Order-Status, Delivery, Stock-Level). The database size depends on the number of warehouses. The workload has many writes and a lot of contention, so both disk speed and locking affect the results.

BenchBase [2, 3] is a benchmarking tool that supports many databases. It reports *throughput* (completed requests per second), *goodput* (successfully committed requests per second) and latency.

Isolation levels. PostgreSQL's default isolation level is `READ COMMITTED` [4]. In `SERIALIZABLE` mode it uses Serializable Snapshot Isolation [5], which tracks reads using `SIRead` (predicate) locks and aborts transactions that could break serializability. `SIRead` locks don't block other transactions, but they take up space in the shared lock table, whose size is fixed at server start [6, 7].

I/O and caching. On paper, the two engines start out equal: both have a small internal cache by default, 128 MiB for PostgreSQL's `shared_buffers` [8] and for MariaDB's `innodb_buffer_pool_size` [9]. The difference is what happens on a cache miss, and this turned out to be the main obstacle to comparing them:

- PostgreSQL uses buffered I/O, so a miss goes through the operating system, which keeps the data in its page cache [8]. The page cache has no fixed limit and grows into all free RAM on the machine, so PostgreSQL effectively gets a much larger second-level cache.
- MariaDB's InnoDB engine uses direct I/O by default [10], bypassing the page cache. It only has its 128 MiB buffer pool, and every miss goes to the disk.

3. Experimental Setup

- **Hardware:** the faculty's student environment, an `x86_64` Debian 13 server with 24 logical CPUs (as reported by MariaDB) and 128 GB of DDR4 RAM, which hosts other students' databases at the same

time. It has one data directory on an SSD and one on a RAID array of HDDs, and both engines used the same two drives.

- **Software:** PostgreSQL 17.8 (before the server restart) and 17.9 (after), MariaDB 11.8.6.
- **Workload:** BenchBase TPC-C, scale factor 1000 (about 70 GB per database), rate set to 100,000 (in practice, no rate limit).
- **Procedure:** I ran the benchmark with 2, 1, 10, 20, 30, 40, 50, 60 and 70 terminals, 20 minutes each, with a 120 s warmup and a 120 s pause between runs. The 2-terminal run went first and only served as the starting point for the database counters, so I don't report it. A Python script generated the BenchBase configs and ran everything.
- **Metrics:** throughput, mean and p95 latency from BenchBase, plus database counters (`pg_stat_database` [11] for PostgreSQL, InnoDB buffer pool statistics for MariaDB), measured as the change during each run.
- **PostgreSQL:** server defaults (`shared_buffers` = 128 MB) except `work_mem` = 128 MB, with READ COMMITTED isolation (except in Section 4.1).
- **MariaDB:** server defaults (`innodb_buffer_pool_size` = 128 MiB) in three variants:
 - a) Plain defaults. MariaDB ignores BenchBase's isolation setting, so these runs used its default, REPEATABLE READ [12].
 - b) READ COMMITTED set manually, `sort_buffer_size` and `join_buffer_size` set to 128 MiB to match PostgreSQL's `work_mem`, and `innodb_io_capacity/_max` raised from 200/2000 to 2000/4000 [10].
 - c) Like (b), but with direct I/O turned off (`innodb_data_file_buffering` and `innodb_log_file_buffering` set to ON; in MariaDB 11.8 these replace `innodb_flush_method` [10]).
- **Durability:** both engines used their safe defaults (`fsync` and `synchronous_commit` on in PostgreSQL, `innodb_flush_log_at_trx_commit` = 1 and the doublewrite buffer in MariaDB).

4. Results

4.1 First round and the SERIALIZABLE problem

In the first round I ran PostgreSQL with SERIALIZABLE isolation and tested the number of terminals, `random_page_cost` (1.0–6.0) and the CPU cost




Figure 1: First round (SERIALIZABLE): throughput and goodput vs. number of terminals, SSD and HDD.

parameters (`cpu_tuple_cost`, `cpu_index_tuple_cost`, `cpu_operator_cost`, scaled together from 1/8x to 16x their defaults) [13]. These runs only kept BenchBase summaries, without database counters.

On SSD, goodput followed throughput closely and reached 2292 tps at 60 terminals (Figure 1). On HDD, throughput grew to 1340 tps, but goodput stayed between 113 and 169 tps at every terminal count from 10 upwards, so most HDD transactions were aborted. With slower I/O, transactions take longer, overlap more, and are more likely to conflict under SSI [5].

Neither cost parameter had a clear effect. On SSD, throughput stayed within about $\pm 5\%$ (872–904 tps at 10 terminals for `random_page_cost`, 2115–2309 tps at 60 terminals for the CPU costs). On HDD, goodput stayed between 116 and 186 tps for all values.

When I came back for a second, more detailed round, I couldn't reproduce the SSD results: the SSD was now even slower than the HDD. The main bottleneck turned out to be locking: SERIALIZABLE mode creates many SIRead locks, and they count towards the shared lock limit [6, 7]. I switched to READ COMMITTED, which got rid of the lock problem, and used it for all later runs.

4.2 PostgreSQL: SSD vs. RAID HDD

Even with READ COMMITTED, the second round was much slower than the first, peaking at about 300 tps on HDD and 400 tps on SSD. Only after the server was restarted did the results return to the level of the first round. Table 1 and Figure 2 show these post-restart results. The configuration was identical before and after the restart, except that the server was updated from PostgreSQL 17.8 to 17.9. I couldn't find out what caused the slowdown, so I only mention it here as a caveat.

Table 1: PostgreSQL, READ COMMITTED, after the server restart. Throughput in transactions per second (tps), latency in ms.

Terminals	SSD tps	HDD tps	SSD/HDD	SSD avg lat.	HDD avg lat.	SSD p95	HDD p95
1	88	63	1.40	11.3	15.8	29.1	45.1
10	721	248	2.90	13.9	40.1	39.2	150.9
20	1225	346	3.54	16.3	57.7	43.8	250.8
30	1685	505	3.34	17.8	59.4	45.6	270.1
40	2092	794	2.63	19.1	50.3	47.7	225.3
50	2310	1225	1.89	21.6	40.8	54.3	159.8
60	2383	1710	1.39	25.2	34.9	62.9	116.8
70	2546	2161	1.18	27.5	32.4	72.9	109.4




Figure 2: PostgreSQL after the server restart: throughput (blue) and average latency (red) vs. number of terminals. Left: SSD. Right: RAID HDD.

Table 2: Best throughput for each configuration (tps, number of terminals in parentheses).

Configuration	SSD	HDD
PostgreSQL (after restart)	2546 (70)	2161 (70)
MariaDB (a): defaults (REPEATABLE READ)	772 (70)	62 (30)
MariaDB (b): READ COMMITTED + tuned buffers and I/O capacity	760 (70)	63 (30)
MariaDB (c): tuned, direct I/O off	860 (20)	—

SSD throughput grows steadily and starts to level off after 50 terminals, while its latency grows slowly with load. HDD throughput grows faster and faster, and at 70 terminals it's only 18% behind the SSD. HDD latency even drops after 30 terminals. My guess is that this is thanks to the RAID array: with more requests waiting at once, more disks can work in parallel. The cache hit ratio at 70 terminals is practically the same on both drives (88.5% on SSD, 88.4% on HDD), so the difference comes from how long a cache miss takes, not how often it happens. The SSD's latency is lower at every load, with p95 latency 1.5–5.7x lower.

4.3 PostgreSQL vs. MariaDB

MariaDB was about 3.3x slower than PostgreSQL on SSD and about 35x slower on HDD. Variant (b), including the switch from REPEATABLE READ to READ COMMITTED, changed nothing. Turning off direct I/O in variant (c) made MariaDB faster with few terminals (728 vs. 463 tps at 10 terminals), but it peaked at 20 terminals and then dropped to 570 tps at 70, below both other variants (Figure 3). On HDD, MariaDB stayed at about 60 tps from 20 terminals onwards, with p95 latency over 3 s at 60–70 terminals.

These numbers should not be read as “PostgreSQL is 3–35x faster than MariaDB”. Both engines had the same small internal cache, but PostgreSQL's misses were often served from the OS page cache, which could grow into all free RAM on the server. MariaDB's misses went straight to the disk, and I couldn't enlarge its buffer pool without a server restart. The server has 128 GB of RAM, so with about 70 GB of data, a large part of PostgreSQL's database could fit in the page cache, while MariaDB could keep only 128 MiB of it in memory. This difference in effective cache size is probably the main reason for the gap, especially on HDD, where every cache miss is expensive. Turning off direct I/O in variant (c) did not make the setups equivalent either, since InnoDB still manages its own buffer pool and throughput got worse at high load. A fair comparison would need both engines to have the same amount of memory available





Figure 3: MariaDB on SSD: throughput (blue) and average latency (red) vs. number of terminals. Left: variant (a), defaults. Right: variant (c), direct I/O off.

for caching.

5. Discussion

- **Is the SSD worth it?** With moderate load the SSD gives about 3x the throughput, but under high load the HDD array gets within about 20%. The clearest advantage of the SSD is its consistently low latency. Whether that’s worth the price depends on the expected load and on how much latency matters.
- **Page cache and the drive comparison.** PostgreSQL’s use of the OS page cache probably hides part of the difference between the drives, because `blks_read` also counts reads served from the page cache rather than the disk [11]. So some “disk reads” in my measurements may never have reached either drive.
- **Page cache and the engine comparison.** Because of the different caching models (Section 2), my PostgreSQL vs. MariaDB results compare cache sizes more than they compare engines. I can’t say from this data which engine is faster. Because the server was shared, I also don’t know how much of its 128 GB was actually free for the page cache during each run, which is another source of noise for PostgreSQL.
- **Tuning mattered less than the environment.** Planner cost parameters and MariaDB’s I/O settings had little or no effect. The biggest changes came from outside the database configuration: aborts and lock exhaustion under `SERIALIZABLE`, and the unexplained slowdown that only went away after the server restart.

6. Limitations

- The server was shared with other students’ databases, and I ran each configuration only once, so the results include noise I couldn’t control. The unexplained difference before and after the server

restart shows this.

- The PostgreSQL vs. MariaDB comparison isn’t valid as an engine comparison. Both had the same small internal cache, but only PostgreSQL could also use the OS page cache, which has no fixed limit. Giving MariaDB a comparable buffer pool requires a server restart I didn’t have permission for.
- With warmup enabled, BenchBase sometimes reports goodput higher than throughput. This is a known issue [14], so I report throughput.

7. Future Work

- Run the tests on a server that isn’t shared with other users.
- Tune buffer sizes for both engines, especially `innodb_buffer_pool_size`, which needs a server restart.
- Repeat each configuration several times to get error bars.

8. Conclusion

On this server, PostgreSQL on SSD was up to about 3.5x faster than on RAID HDD with moderate load, but only about 1.2x faster with 70 terminals. Lower latency was the SSD’s consistent advantage. MariaDB was 3–35x slower than PostgreSQL, but because PostgreSQL could use the OS page cache and MariaDB could not, this reflects the difference in available cache rather than in the engines themselves.

References

- [1] Transaction Processing Performance Council. *TPC Benchmark C, Standard Specification*, Revision 5.11, 2010. <https://www.tpc.org/tpcc/>
- [2] D. E. Difallah, A. Pavlo, C. Curino, P. Cudré-Mauroux. “OLTP-Bench: An Extensible Testbed for

Benchmarking Relational Databases.” *PVLDB* 7(4), 2013.

[3] BenchBase. <https://github.com/cmu-db/benchbase>

[4] PostgreSQL 17 Documentation. “Transaction Isolation.” <https://www.postgresql.org/docs/17/transaction-iso.html>

[5] D. R. K. Ports, K. Grittnner. “Serializable Snapshot Isolation in PostgreSQL.” *PVLDB* 5(12), 2012.

[6] PostgreSQL 17 Documentation. “Lock Management” (`max_locks_per_transaction`, `max_pred_locks_per_transaction`). <https://www.postgresql.org/docs/17/runtime-config-locks.html>

[7] “Post-mortem: Postgres out of shared memory error.” <https://baida.dev/articles/post-mortem-postgres-out-of-shared-memory-error>

[8] PostgreSQL 17 Documentation. “Resource Consumption” (`shared_buffers`, `work_mem`). <https://www.postgresql.org/docs/17/runtime-config-resource.html>

[9] MariaDB Documentation. “InnoDB Buffer Pool.” <https://mariadb.com/docs/server/server-usage/storage-engines/innodb/innodb-buffer-pool>

[10] MariaDB Documentation. “InnoDB System Variables.” <https://mariadb.com/docs/server/server-usage/storage-engines/innodb/innodb-system-variables>

[11] PostgreSQL 17 Documentation. “The Cumulative Statistics System” (`pg_stat_database`). <https://www.postgresql.org/docs/17/monitoring-stats.html>

[12] MariaDB Documentation. “SET TRANSACTION” (isolation levels). <https://mariadb.com/docs/server/reference/sql-statements/administrative-sql-statements/set-commands/set-transaction>

[13] PostgreSQL 17 Documentation. “Query Planning: Planner Cost Constants.” <https://www.postgresql.org/docs/17/runtime-config-query.html>

[14] BenchBase issue #606. <https://github.com/cmu-db/benchbase/issues/606>